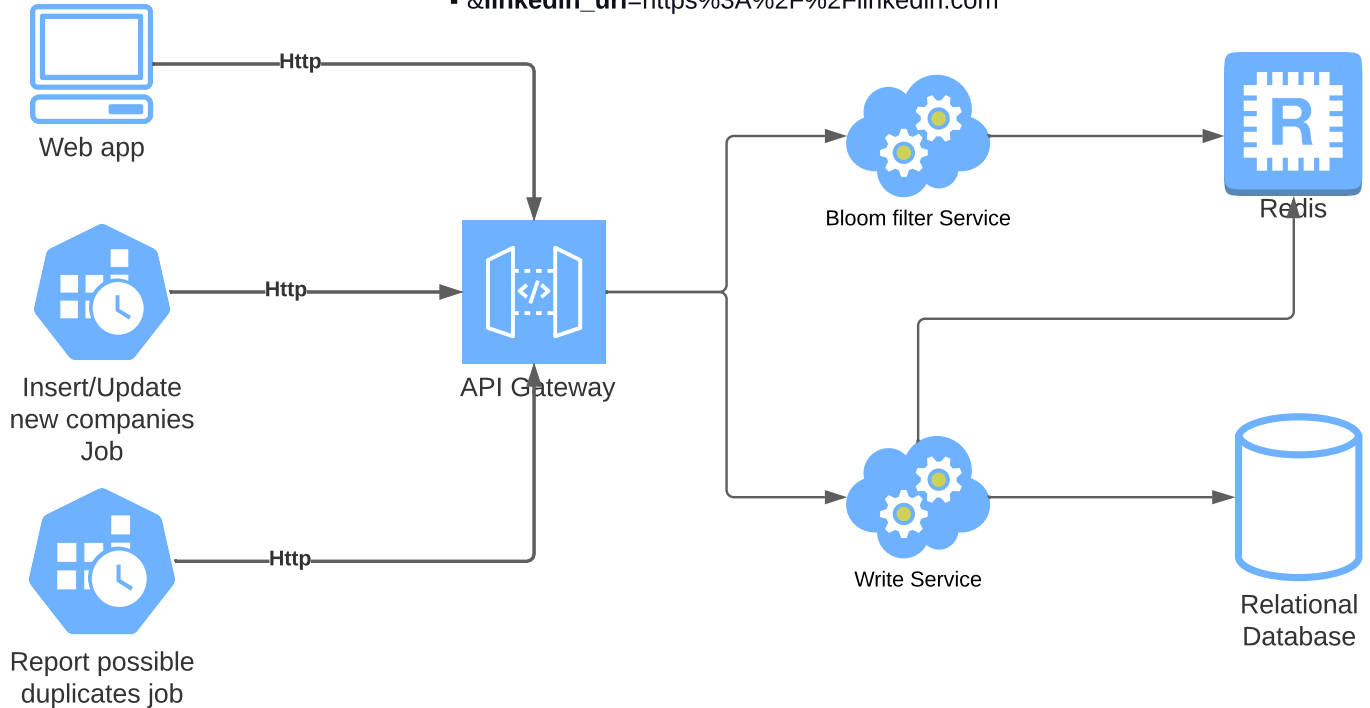


### API Endpoints

- PATCH /api/companies/:companyId
- POST /api/companies
- GET /api/companies
  - ?name=company1
  - &short\_name=c1
  - &website\_url=https%3A%2F%2Fgoogle.com
  - &linkedin\_url=https%3A%2F%2Flinkedin.com



- I have considered switching the database solution to a NoSql database but, as consistency tends to be important for company information, the relational database sounds like the best solution;
- One of the tradeoffs of the Redis cache would be data synchronization between the cache and the relational database;
- The "report possible duplicates job" would be executed by a scheduler and would generate a report with the possible duplicates;
- The Bloom filter is probabilistic so false positives are expected.